

Akhil Bodahanapati

akb198@pitt.edu | 747-312-1646 | akhilbod.github.io | linkedin.com/in/akhil-bod | github.com/akhilBod

EDUCATION

University of Pittsburgh

Expected April 2028

Bachelor of Science in Computer Science

- Relevant Coursework: Data Structures, Algorithms, Operating Systems, Compilers, Computer Networks, Systems Software, Probability and Statistics, Multivariable Calculus

EXPERIENCE

Nutanix

May 2026 – August 2026

Incoming Software Engineer Intern

San Jose, CA

- Building control plane APIs in Python and C++ to automate VM lifecycle management

Natural Language Processing and AI Laboratory (NAIL)

December 2025 – April 2026

Software Engineer Intern

Pittsburgh, PA

- Built async Python APIs with OpenAPI contracts powering NLP inference pipelines over 1M+ documents
- Standardized integration with downstream research services through versioned definitions and async architecture
- Re-architected batch job processing onto a Redis task queue with priority routing, cutting average job wait time from (4 min → 75 sec) and tripling pipeline throughput

Federated Hermes

May 2025 – August 2025

Software Engineer Intern

Pittsburgh, PA

- Designed and shipped a Python microservice for portfolio rebalancing across 2000+ users and \$2B+ AUM
- Built an observability platform surfacing live p99 latency and error rates across 8 services, cutting MTTR 95%
- Automated a 12-step manual deployment pipeline into a single idempotent CLI script with dry-run mode and rollback support, cutting release time by 65% and eliminating a recurring class of partial-deploy errors

Star Hospitals

August 2024 – Present

Software Engineer Intern

Los Angeles, CA

- Led a 3-person team to design and ship a patient management system handling 1,200+ daily transactions
- Built check-in, triage severity scoring, and discharge APIs as independently deployable services, cutting average patient wait time 38% (45 → 28 min) and surfacing real-time floor capacity to clinical staff

PROJECTS

InterviewSense | Python, React, PostgreSQL, Redis, AWS, Docker

- Built an interview prep platform serving 10K+ users, with real-time feedback across multiple interview types
- Built an async job pipeline that streams LLM evaluation responses to maintain sub-200ms API response times
- Designed a PostgreSQL with Redis caching, with sub-120ms p99 read latency at 300+ concurrent sessions

Quantitative Investment Engine | Python, SQL, Databricks, AWS

- Won 1st place in a company-wide hackathon; selected for Q3 2025 production with \$75K annual cost savings
- Built a research engine for the portfolio construction team, ingesting live market data, ran quantitative factor models, and surfaced AI due diligence reports to serve as a third-party perspective on portfolio

Orbit: Distributed Vector Search Engine | Python, C++, FAISS, AWS, Docker, Redis

- Won Most Significant Innovation among 700+ participants at CMU TartanHacks 2024
- Built a tri-modal search backend combining FAISS dense retrieval and reranking over 500K+ records
- Achieved sub-100ms query latency with optimized C++ bindings and pipeline parallelism across retrieval stages

TECHNICAL SKILLS

Languages: Python, C++, Go, Java, TypeScript, SQL, Bash

Frameworks: Flask, Spring Boot, React, Celery, PySpark

Infrastructure: AWS (ECS, Lambda, SQS, S3, RDS), Docker, Kubernetes

Data Stores: PostgreSQL, Redis, Databricks

Tools: Git, GitHub Actions, CI/CD, Linux, Prometheus